

ISHAN SHAH

647-765-0933 · i9shah@uwaterloo.ca · LinkedIn [↗](#) · GitHub [↗](#) · ishanshah.org [↗](#)

EDUCATION

University of Waterloo
Bachelor of Computer Science

Sept. 2025 – Apr. 2030

EXPERIENCE

Axibo

Research Engineer

Incoming S26

Cambridge, ON

- Developing and deploying Vision-Language-Action (VLA) models and Reinforcement Learning (RL) frameworks to enhance autonomous bimanual manipulation and locomotion capabilities for humanoid robotic platforms.

WAT.ai

Machine Learning Researcher (Robotics)

Jan. 2026 - Present

Waterloo, ON

- Researched **preference learning** and **human-feedback-driven optimization** for sequential decision-making models.
- Designed and implemented **realistic robotic simulation environments** and trained **Actor-Critic (ACT) policies**, evaluating robustness across diverse scenarios.
- Built and optimized **Python/PyTorch learning pipelines** and scalable human-in-the-loop setups for improving policy performance under sparse or under-specified rewards.

IPMD

Software Engineer Intern

July 2025 - Sept. 2025

San Mateo, CA

- Replaced custom auth logic with **Flask-JWT-Extended**, centralizing token management across **10 endpoints** and resolving registration edge cases.
- Modularized frontend by extracting **5+ reusable components**, reducing code duplication by **30%** and accelerating future feature development.
- Built a real-time **computer vision-driven** emotion analysis pipeline integrated into a React (TypeScript) chat interface, surfacing live sentiment feedback during conversations.

TECHNICAL SKILLS

Languages: Python, C++

Frameworks: PyTorch, JAX, ROS2, OpenCV, NumPy

Simulation: MuJoCo, Isaac Lab, Gymnasium

Tools/Deployment: Linux, Git, Docker

PROJECTS

Human-in-the-Loop Teleoperation for Imitation Learning [↗](#) | *MuJoCo, MediaPipe, Pytorch*

- Designed a **low-latency teleoperation** system to control a simulated **ALOHA bimanual robot** using monocular webcam-based hand tracking, eliminating specialized hardware.
- Built a real-time human-to-robot control pipeline using **Inverse Kinematics** to map tracked hand poses into robot joint-space commands, refining **quaternion-based** mappings to avoid gimbal lock.
- Developing a scalable **HDF5-based** demonstration logging pipeline structured for training **Action-Conditioned Transformers (ACT)** and **Diffusion Policy** models.

Reinforcement Learning for Continuous Control [↗](#) | *Python, PyTorch, Gymnasium, Numpy*

- Developed a from-scratch implementation of **Proximal Policy Optimization (PPO)** with **Generalized Advantage Estimation (GAE)** to solve continuous action space environments.
- Architected a custom **Actor-Critic** network in **PyTorch**, optimizing hyperparameters to achieve convergence on the BipedalWalker-v3 environment.
- Implemented **orthogonal weight initialization** and **learning rate annealing**, improving training stability and reducing variance in episodic returns.